

RESEARCH REPORT · JUNE 2026

RemediAX Reporter — Competitive

How RemediAX's HTML Security Reporter compares against
Promptfoo, Garak, PyRIT, Lakera Guard, and Mindgard

Author: Nileshwari Kadgale · Project: RemediAX AI Security Platform · Date: 2026-06-09

GitHub: github.com/nileshvary/nileshvary-ai-security-engine · Live: remediox.streamlit.app

Section 1

Executive Summary

VERDICT

RemediAX ReporterAgent is the most feature-complete free AI security reporter available. It surpasses every open-source competitor and matches or exceeds paid commercial tools on per-finding detail, AI narrative quality, and coverage breadth. Five specific capabilities exist in RemediAX that are not present in any other tool in this analysis.

The AI security tooling landscape has a clear reporting gap: scanners are mature (Garak, PyRIT), but their output is raw machine-readable data unsuitable for security managers, developers, or board-level stakeholders. The only tools producing professional human-readable reports are commercial SaaS products (Mindgard, Lakera) that cost thousands of dollars per year and require vendor lock-in.

RemediAX closes this gap entirely. Its **Agent 3 Reporter** — built on Claude AI + Jinja2 — produces a professional dark-themed HTML report that rivals enterprise commercial tools, includes Claude-generated per-finding analysis that no competitor offers for free, and covers all 20 OWASP vulnerability categories (LLM Top 10 + Agentic Top 10) that no competitor covers at all.

Section 2

RemediAX Reporter Overview

RemediAX Agent 3 (`agents/reporter_agent.py`) is the third stage of the six-agent pipeline: Scanner → Remediator → **Reporter** → Verifier → Orchestrator → CVE Watcher.

Dark-Themed Professional HTML: GitHub-style dark theme rendered via Jinja2 (`templates/summary.html.j2`). Professional for enterprise security reviews.

AI-Generated Executive Summary: Claude (claude-haiku-4-5-20251001) writes a scan-level summary per target. Falls back to deterministic text without an API key.

Per-Finding Cards with Full Context: Each card: attack prompt excerpt, model response excerpt, Claude danger explanation, fix explanation, and guardrail YAML snippet.

OWASP LLM Top 10 Coverage Table (+ ASI01–ASI10): All 10 LLM categories shown with hit/miss and counts. Unique to RemediAX: also covers all 10 ASI Agentic categories in the scanner.

Strategy Breakdown Bar: Visual breakdown of HARDEN / SANITIZE / GUARDRAIL / LOG_ONLY — RemediAX's unique 4-category remediation taxonomy.

17 Unit Tests: `tests/agents/test_reporter_agent.py` — HTML structure, DI paths, Claude API integration, fallback behavior.

Full Competitor Comparison Matrix

Feature	RemediAX	Promptfoo OSS	Garak	PyRIT	Lakera	Mindgard
Report format	HTML (dark-themed)	HTML+JSON+SARIF	JSONL only	SQLite/Azure	SaaS only	SaaS UI+SIEM
Per-finding cards	YES — full context	Partial	NO	NO	NO	YES (paid)
AI narrative/finding	YES — FREE	NO	NO	NO	NO	YES (paid \$\$\$)
AI executive summary	YES — FREE	NO (OSS)	NO	NO	NO	YES (paid)
OWASP LLM Top 10	YES — all 10	YES	LLM01+05 only	Partial	NO	YES+MITRE+NIST
ASI Agentic Top 10	YES — all 10	NO	NO	NO	NO	NO
Guardrail YAML inline	YES — copy-pasteable	NO	NO	NO	NO	NO
Remediation taxonomy	HARDEN/SANITIZE/GUARDRAIL/LOG_ONLY	NO	NO	NO	NO	NO
Severity badges	YES	YES	NO	Partial	NO	YES
Boardroom-ready	YES — FREE	NO (OSS)	NO	NO	NO	YES (paid)
Self-hosted/offline	YES	YES	YES	Partial	NO	NO
Cost	FREE	FREE	FREE	FREE	Paid SaaS	Paid \$\$\$

Competitor Deep Dive

Promptfoo [Closest Open-Source Competitor]

Overview: Promptfoo generates HTML reports with OWASP category mapping, severity levels, and CI/CD integration (SARIF). Its red-team output includes per-test results with attack prompts.

Strengths: OSS tier produces readable HTML with OWASP mapping and JSON/CSV/SARIF export. Strong CI/CD tooling and community adoption.

Gaps vs. Remedix: No AI-generated per-finding explanation. No guardrail YAML. No ASI coverage. Executive summary and PDF export locked behind paid Enterprise tier.

Garak (NVIDIA) [CLI Tool — No Human Report]

Overview: Mature LLM security probe runner with 50+ probes. Primary output is JSONL (raw machine log). HTML digest shows only pass/fail ratios.

Strengths: Offline, comprehensive probe library, NVIDIA-backed. Remedix uses Garak as its scanner engine (tools/garak_runner.py).

Gaps vs. Remedix: No per-finding cards, severity badges, narrative, remediation guidance, or executive summary. Security managers cannot use Garak output directly.

PyRIT (Microsoft) [Azure-Dependent]

Overview: Multi-turn red-teaming framework storing results in SQLite. Azure Foundry UI provides analytics views.

Strengths: Multi-turn crescendo attacks produce realistic findings. Remedix uses PyRIT as its multi-turn attack engine.

Gaps vs. Remedix: No standalone HTML/PDF report. Requires Azure Foundry — not truly self-hosted. No AI explanations or guardrail output.

Lakera Guard [Real-Time SaaS Only]

Overview: Real-time prompt injection detection with SaaS analytics dashboard and SIEM export. Designed for production monitoring, not point-in-time audits.

Strengths: Low latency, SIEM integration, enterprise-grade infrastructure.

Gaps vs. Remedix: No downloadable scan report. SaaS subscription required. Covers only live traffic, not systematic OWASP penetration testing.

Mindgard [Best Commercial Competitor]

Overview: Feature-rich commercial AI red-teaming platform. Maps findings to OWASP + MITRE ATLAS + EU AI Act + NIST AI RMF. Enterprise PDF reports with per-finding attacker context.

Strengths: Richest framework alignment. Boardroom-ready PDFs. Best commercial tool available.

Gaps vs. Remedix: Closed-source, paid, requires sales engagement. No self-hosted option. Missing: ASI Agentic Top 10, guardrail YAML, HARDEN/SANITIZE/GUARDRAIL/LOG_ONLY taxonomy.

RemediAX delivers comparable quality FREE and open-source.

Section 5

5 Unique Differentiators

The following capabilities exist in RemediAX and are **absent from every competitor**:

- 1

AI-Generated Guardrail YAML per Finding — Free

RemediAX is the only tool that generates copy-pasteable deployment-ready guardrail YAML for each specific finding. Claude analyzes the exact attack prompt and produces the YAML inline in the report. No competitor generates deployment-ready guardrail code as report output.
- 2

OWASP ASI Agentic Top 10 (ASI01-ASI10) — All 10 Categories

RemediAX covers all 20 vulnerability categories: OWASP LLM Top 10 (LLM01-LLM10) plus the OWASP Agentic Top 10 (ASI01-ASI10). No competitor covers the ASI framework. As AI agents become the dominant deployment pattern, this is a critical competitive differentiator.
- 3

HARDEN / SANITIZE / GUARDRAIL / LOG_ONLY Remediation Taxonomy

RemediAX maps each OWASP category to a precise named remediation strategy. HARDEN = system prompt patching (LLM01/07), SANITIZE = output redaction (LLM02/05/06), GUARDRAIL = rate-limit/filter rules (LLM10), LOG_ONLY = infrastructure-level issues. No competitor has a named remediation strategy taxonomy — they show the vulnerability but not what class of fix to apply.
- 4

Full Pipeline Free in One Self-Hosted Tool

Scan (Agent 1) + Remediate (Agent 2) + Report (Agent 3) + Verify (Agent 4) + CVE Watch (Agent 6). No other free tool covers the complete AI security lifecycle in one pipeline. Promptfoo covers scan+report only. Garak/PyRIT cover scan only. No competitor covers remediation + report + verification free.
- 5

Claude AI Context Per Finding — Free, No Subscription Required

Claude-generated danger explanations, fix text, and guardrail YAML per finding — available free with any API key. Mindgard provides AI context as a paid SaaS feature only. RemediAX delivers the same quality to individuals and small teams with no budget. Report is fully functional without any API key using pre-written OWASP fallback content.

Section 6

Conclusion

MARKE
T POSIT
ION

RemediAX is the best free AI security reporter available. Among open-source tools it is unambiguously ahead — Promptfoo OSS is the only real competitor and it lacks AI narrative, guardrail YAML, ASI coverage, and remediation strategy taxonomy. Among commercial tools, RemediAX matches Mindgard's per-finding quality for free, and exceeds Mindgard on ASI coverage, guardrail YAML, and open-source accessibility.

Where RemediAX Can Improve: PDF export (Mindgard and Promptfoo Enterprise produce native PDFs). CI/CD-native SARIF output would accelerate enterprise adoption — Agent 4 (Verifier) addresses this.

Strategic Recommendation: Position RemediAX explicitly as "the free Mindgard alternative for open-source AI teams." The gap between Mindgard (closed, expensive) and the open-source field (Garak/PyRIT: raw data, no reports) is exactly where RemediAX competes. The five unique differentiators — especially ASI coverage and free AI-generated guardrails — are authentic advantages to communicate in the README, HackerOne report, and conference submissions.

Section 7

Sources

- Promptfoo GitHub — github.com/promptfoo/promptfoo
- Promptfoo Red Team Docs — promptfoo.dev/docs/red-team/
- Promptfoo Enterprise Findings — promptfoo.dev/docs/enterprise/findings/
- Promptfoo Remediation Reports (Enterprise) — promptfoo.dev/docs/enterprise/remediation-reports/
- Garak NVIDIA GitHub — github.com/NVIDIA/garak
- Garak CLI Reference — reference.garak.ai/en/latest/cliref.html
- NVIDIA Garak Review 2026 — appsecsanta.com/garak
- PyRIT Documentation — microsoft.github.io/PyRIT/
- Microsoft AI Red Teaming Agent — devblogs.microsoft.com/foundry/ai-red-teaming-agent-preview/
- PyRIT 2026 Review — appsecsanta.com/pyrit
- Lakera Guard Documentation — docs.lakera.ai/guard
- Mindgard Automated AI Red Teaming — mindgard.ai/automated-ai-red-teaming
- Mindgard 2026 Review — appsecsanta.com/mindgard
- OWASP LLM Top 10 Project — owasp.org/www-project-top-10-for-large-language-model-applications/
- OWASP Agentic Top 10 2026 — trydeepteam.com/docs/frameworks-owasp-top-10-for-agentic-applications